

AI Security Lab — Level 06: Bypass a Refusal

Vulnerability Overview

- **Vulnerability Type:** Refusal Mechanism Evasion / Safety-Control Robustness
- **Objective:** Evaluate the robustness of an AI model's refusal mechanisms against variations in prompt structure, semantic framing, and contextual manipulation.

Technical Analysis

Modern AI safety systems may use combinations of classifiers, policy models, semantic analysis, instruction hierarchy, and model-level alignment to identify and refuse disallowed requests.

However, security controls that depend primarily on simple pattern matching or keyword detection may be vulnerable to changes in linguistic structure and contextual framing.

This laboratory evaluates whether a refusal mechanism remains effective when the underlying intent is expressed through different forms of language rather than relying on a single recognizable keyword or phrase.

Security Impact

Potential impacts include:

- Inconsistent enforcement of safety policies.
- Evasion of simplistic keyword-based filters.
- Increased exposure to adversarial prompt techniques.
- Generation of responses outside the application's intended safety boundaries.

Remediation

- Implement **multi-layered semantic classification** and intent-based moderation.
- Combine model-level safeguards with application-level security controls.
- Evaluate security controls against paraphrased and adversarial prompt variations.
- Maintain continuous red-team testing to identify newly emerging evasion techniques.
- Ensure high-risk actions require independent authorization controls rather than relying solely on model refusal behavior.

Author: Moe Htet Ar Kar (Phoe Cho)